

WESTCLIFF UNIVERSITY
College of Technology & Engineering
In Partnership with Presidential Graduate School (Nepal)

TECH 400
Introduction to Information Retrieval
Comprehensive Lecture Notes

Theory, Mathematics, and Worked Exercises for Every Topic

Aligned to the TECH 400 course outline (Weeks 1–8) and to
Manning, C., Raghavan, P., & Schütze, H. (2009). *An Introduction to Information
Retrieval*. Cambridge University Press.
(Freely available at <http://www.informationretrieval.org/>)

Contents

How to Use These Notes	4
Week 1: Boolean Retrieval and the Term Vocabulary	5
1 What Is Information Retrieval?	5
2 The Term–Document Incidence Matrix	5
3 The Boolean Retrieval Model	6
3.1 The Inverted Index	6
3.2 Processing Boolean Queries: The Merge Algorithm	7
3.3 Optimizing Multi-Term AND Queries	8
4 The Term Vocabulary	8
4.1 Tokenization	8
4.2 Stop Words	9
4.3 Normalization and Case Folding	9
4.4 Stemming and Lemmatization	9
Week 2: Dictionaries, Tolerant Retrieval, and Index Construction	11
5 Dictionary Data Structures	11
6 Wildcard Queries and the Permuterm Index	11
7 Spelling Correction	12
7.1 Edit Distance	12
7.2 Jaccard Coefficient for k -gram Candidate Generation	13
8 Index Construction	13
8.1 Blocked Sort-Based Indexing (BSBI)	13
8.2 Single-Pass In-Memory Indexing (SPIMI)	14
8.3 Distributed and Dynamic Indexing (Conceptual Overview)	15
Week 3: Term Weighting, the Vector Space Model, and Efficient Scoring	16
9 Parametric and Zone Indexes	16
10 Term Frequency and Document Frequency	16
10.1 Inverse Document Frequency	16
10.2 tf-idf Weighting	17
11 The Vector Space Model	17
12 Efficient Scoring	19
12.1 Weighted Zone Scoring	19
12.2 The BM25 Scoring Function	19

Week 4: Evaluation, Relevance Feedback, and Query Expansion	21
13 Evaluating an Unranked Retrieval Set	21
14 Evaluating a Ranked Retrieval List	21
14.1 Average Precision and Mean Average Precision	22
14.2 Normalized Discounted Cumulative Gain (nDCG)	22
15 Relevance Feedback: The Rocchio Algorithm	24
16 Probabilistic Relevance Feedback (Preview)	24
17 Query Expansion	25
Week 5: Probabilistic Information Retrieval and Language Models for IR	26
18 The Probability Ranking Principle	26
19 The Binary Independence Model (BIM)	26
20 Probabilistic Relevance Feedback	27
21 Language Models for Information Retrieval	28
21.1 Jelinek–Mercer Smoothing	28
21.2 Dirichlet Smoothing	29
Week 6: Text Classification and Naïve Bayes	30
22 The Text Classification Problem	30
23 The Naive Bayes Generative Model	30
24 Evaluating a Classifier	32
Week 7: Web Search Basics and Link Analysis	33
25 Characteristics of the Web as a Retrieval Problem	33
26 Web Crawling (Conceptual Overview)	33
26.1 Estimating Relative Search-Engine Coverage from Overlap	33
27 Link Analysis: PageRank	34
27.1 HITS: Hubs and Authorities (Brief)	36
Week 8: Course Review and Synthesis	37
28 The Big Picture: How the Models Connect	37
29 Master Formula Sheet	38

30 Comprehensive Review Problems	38
Appendix	41
Further Reading	41
Notation Glossary	41

How to Use These Notes

These notes follow the TECH 400 course outline week by week. Each week contains three kinds of material, visually distinguished by colored boxes:

- **Definition / Theory** boxes state the core concepts precisely.
- **Key Formula** boxes isolate the mathematics you are responsible for – these are the equations you should be able to reproduce and apply from memory.
- **Worked Example** boxes show a complete, fully computed example, step by step, using small numbers chosen so the arithmetic can be done with pencil, paper, and (at most) a basic calculator for logarithms and square roots.
- **Practice Exercise** boxes give you a very similar problem to solve yourself. A full **Solution** is provided immediately afterward so you can check your work – cover it with a sheet of paper while you attempt the problem, and only look once you have an answer.

Every numerical example in this document has been verified by hand; where a base is chosen for logarithms (base 2 vs. base 10), it is stated explicitly, because *the base only rescales every weight by the same constant and never changes which documents rank higher* – this is proven in Section 10.2.

Week 1

Boolean Retrieval and the Term Vocabulary

1 What Is Information Retrieval?

Definition 1.1 – Information Retrieval

Information Retrieval (IR) is the activity of finding material – usually documents – of an unstructured nature (typically text) that satisfies an information need, from within a large collection (a *corpus*).

IR differs from a database (DB) query in a fundamental way:

	Database Query	IR Query
Data	Structured (rows/columns, fixed schema)	Unstructured (free text)
Semantics	Well defined (SQL is exact)	Ambiguous; keywords are a proxy for meaning
Result	Exact match, set of rows	<i>Ranked list</i> , degree of relevance
Answer quality	Right or wrong	Better or worse (relevance is a matter of degree)

The central object of study is the **collection** (or **corpus**): a fixed set of N documents $\{d_1, \dots, d_N\}$ over which a **query** q is evaluated to produce a ranked or unranked subset.

2 The Term–Document Incidence Matrix

Suppose we restrict attention to a small, fixed set of terms of interest. For each term t and document d we can ask a single yes/no question: *does t occur in d ?* Recording the answer for every (term, document) pair gives the **term–document incidence matrix**.

Worked Example 1.1 – Building an Incidence Matrix

Consider a five-document toy corpus about programming:

Doc	Text
d_1	“Python is a scripting language used for code automation”
d_2	“Java is a compiled language, not a scripting language”
d_3	“Python and Java are both popular for writing code”
d_4	“Scripting languages are easy to learn”
d_5	“Compiled code runs faster than scripted code”

Restricting to the four terms {python, java, code, easy} (using exact word forms, before any stemming is applied – stemming is introduced in Section 4.4), the incidence matrix is:

Term	d_1	d_2	d_3	d_4	d_5
python	1	0	1	0	0
java	0	1	1	0	0
code	1	0	1	0	1
easy	0	0	0	1	0

A Boolean query is answered by taking the appropriate bitwise combination of rows. For example, **python AND code** combines row **python** = (1, 0, 1, 0, 0) with row **code** = (1, 0, 1, 0, 1) using bitwise AND, giving (1, 0, 1, 0, 0) – documents d_1 and d_3 .

Remark: An incidence matrix is simple but wasteful: it is $|V| \times N$ (vocabulary size times number of documents) and is almost entirely zeros for a real collection (most words do not appear in most documents). This motivates the inverted index of Section 3.1.

3 The Boolean Retrieval Model

Definition 1.2 – Boolean Retrieval Model

The **Boolean retrieval model** answers a query that is a Boolean expression of terms combined with the operators AND, OR, and NOT. A document either *satisfies* the expression or it does not – there is no notion of degree or ranking.

Boolean retrieval was the dominant commercial IR model for decades (library catalogs, legal/-patent search) because it gives the user precise, predictable control, at the cost of requiring the user to construct the query expression correctly and offering no ranking of the results.

3.1 The Inverted Index

Definition 1.3 – Inverted Index

An **inverted index** consists of:

1. a **dictionary** (vocabulary) of all terms in the collection, and
2. for each term t , a **postings list**: a sorted list of the document IDs (docIDs) of every document containing t (optionally annotated with term frequency and positions).




Figure 1.1: Inverted index for the corpus in Worked Example 1.1 (dictionary → sorted postings lists of docIDs).

3.2 Processing Boolean Queries: The Merge Algorithm

To answer AND and OR queries efficiently, we exploit the fact that postings lists are kept *sorted* by docID. Two sorted lists can be intersected or unioned in a single linear pass using two pointers.

Key Formula 1.1 – Cost of the Merge Algorithm

If two postings lists have lengths x and y , intersecting (or unioning) them by the merge algorithm takes

$$O(x + y)$$

comparisons in the worst case – each pointer only ever advances forward, and every advance is paid for by one comparison.

Algorithm 1 INTERSECTPOSTINGS(p_1, p_2)

```

1: answer  $\leftarrow \langle \rangle$ 
2: while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$  do
3:   if  $\text{docID}(p_1) = \text{docID}(p_2)$  then
4:     APPEND(answer,  $\text{docID}(p_1)$ )
5:      $p_1 \leftarrow \text{next}(p_1)$ ;  $p_2 \leftarrow \text{next}(p_2)$ 
6:   else if  $\text{docID}(p_1) < \text{docID}(p_2)$  then
7:      $p_1 \leftarrow \text{next}(p_1)$ 
8:   else
9:      $p_2 \leftarrow \text{next}(p_2)$ 
10:  end if
11: end while
12: return answer

```

Worked Example 1.2 – Merging Postings by Hand

Using Figure 1.1, evaluate python AND code: postings are $\langle 1, 3 \rangle$ and $\langle 1, 3, 5 \rangle$.

Step	p_1 points at	p_2 points at	Compare	Action
1	1	1	$1 = 1$	output 1; advance both
2	3	3	$3 = 3$	output 3; advance both
3	NIL	5	–	p_1 exhausted, stop

Result: $\langle 1, 3 \rangle$, found in 2 comparisons (here $x = 2, y = 3$, and $2 \leq x + y = 5$, consistent with Key Formula 1.1).

Now evaluate code AND NOT python: start from $\langle 1, 3, 5 \rangle$ (code) and remove any docID that also appears in $\langle 1, 3 \rangle$ (python) by the same merge logic, keeping only the elements of the first list with no match in the second: 1 matches (drop), 3 matches (drop), 5 has no match in python's list (keep). Result: $\langle 5 \rangle$.

Practice Exercise 1.1

Using the postings lists of Figure 1.1 (python = $\langle 1, 3 \rangle$, java = $\langle 2, 3 \rangle$, code = $\langle 1, 3, 5 \rangle$, easy = $\langle 4 \rangle$):

- Compute java OR easy by merging, showing your pointer-walk table.
- Compute python AND java by merging.
- How many total comparisons did the merge in part (b) take? Confirm it does not exceed $x + y$.

Solution. (a) Merging $\langle 2, 3 \rangle$ and $\langle 4 \rangle$ (union: keep every element, output the smaller and advance, output both and advance both on a tie): compare 2 vs 4 $\Rightarrow$ output 2, advance java pointer; compare 3 vs 4 $\Rightarrow$ output 3, advance java pointer; java list exhausted, flush remainder of easy list: output 4. Result: $\langle 2, 3, 4 \rangle$.

(b) Merging $\langle 1, 3 \rangle$ and $\langle 2, 3 \rangle$: compare 1 vs 2 $\Rightarrow$ no match, advance smaller (python); compare 3 vs 2 $\Rightarrow$ no match, advance smaller (java); compare 3 vs 3 $\Rightarrow$ match, output 3, advance both; both lists exhausted. Result: $\langle 3 \rangle$.

(c) 3 comparisons were used; $x + y = 2 + 2 = 4 \geq 3$. $\checkmark$

3.3 Optimizing Multi-Term AND Queries

When a query has more than two AND-ed terms, the *order* in which postings lists are merged matters a great deal, because each intermediate result becomes the input to the next merge.

Key Formula 1.2 – Greedy Merge Order

For a query t_1 AND t_2 AND $\dots$ AND t_k , process the terms in order of **increasing document frequency** df_t (equivalently, increasing postings-list length). Always merge the current (shrinking) intermediate result with the next-shortest remaining list. This minimizes the total work because the intermediate result can never grow larger than the shortest list merged so far.

Practice Exercise 1.2

A query engine reports the following document frequencies: $df(\text{database}) = 500$, $df(\text{index}) = 1,000$, $df(\text{query}) = 50$. For the query **database AND index AND query**, in what order should the three postings lists be merged? Explain using Key Formula 1.2.

Solution. Sort terms by increasing df : **query** (50) < **database** (500) < **index** (1,000). So first intersect **query** and **database** (the two shortest lists), producing an intermediate result no longer than 50 items; then intersect that result with **index**. Processing **index** first (length 1,000) against another long list would waste comparisons on many docIDs that are doomed to be eliminated later.

4 The Term Vocabulary

Before any index can be built, raw text must be converted into a stream of index terms. This pipeline is called **tokenization and normalization**.

4.1 Tokenization

Definition 1.4 – Token, Type, Term

A **token** is an instance of a sequence of characters in a document that is grouped together as a useful semantic unit (e.g. one occurrence of “code”). A **type** is the class of all tokens with the same character sequence. A **term** is a (possibly normalized) type that is entered into the IR system’s dictionary.

Tokenization raises many non-trivial decisions: should Berkeley’s split into Berkeley + ’s? Should co-education be one token or two? Is 192.168.1.1 one token? Real systems resolve these with language- and application-specific rules; for this course, assume simple whitespace/punctuation splitting unless told otherwise.

4.2 Stop Words

Definition 1.5 – Stop Words

Stop words are extremely common words (*the, a, is, of, and, . . .*) that are often *excluded* from the index because they have very low discriminating power and occur in nearly every document (recall Worked Example 3.1 later, where a term occurring in *every* document receives an *idf* of exactly 0).

Modern web search engines tend to keep stop words (indexed cheaply, e.g. via impact-ordered postings) because they matter for phrase queries (“to be or not to be”), but the classical exclusion technique remains important for compact indexes and for simpler classroom examples.

4.3 Normalization and Case Folding

Normalization maps superficially different tokens that should be treated as the same term into one canonical form: U.S.A. → USA, anti-discriminatory → antidiscriminatory. **Case folding** reduces all letters to lower case so that Automobile and automobile become the same term (with care taken for mid-sentence proper nouns and acronyms, e.g. General Motors vs. gm).

4.4 Stemming and Lemmatization

Definition 1.6 – Stemming vs. Lemmatization

Stemming crudely chops off word endings using heuristic rules to conflate related forms (automate, automatic, automation → automat). **Lemmatization** uses vocabulary and morphological analysis to return the dictionary base form, or *lemma* (is, was, be → be).

The Porter stemmer is the classic English stemming algorithm, applying an ordered sequence of suffix-stripping rules (e.g. SSES → SS, IES→I, ATIONAL→ATE). For hand exercises we use a small simplified rule table capturing the flavor of Porter’s algorithm without its full complexity:

Rule	Condition	Suffix	Replacement
S1	word ends in	ies	→ y
S2	word ends in	sses	→ ss
S3	word ends in	ing (stem has a vowel)	→ (drop)
S4	word ends in	ed (stem has a vowel)	→ (drop)
S5	word ends in	s (not ss)	→ (drop)

Worked Example 1.3 – Tokenize, Remove Stop Words, Stem

Sentence: “The runners were running quickly and easily across the fields.”

Step 1 (tokenize & case-fold): the, runners, were, running, quickly, and, easily, across, the, fields

Step 2 (remove stop words) using stoplist {the, were, and}: runners, running, quickly, easily, across, fields

Step 3 (apply stemming rules S1–S5):

Token	Rule applied	Stem
runners	S5 (s dropped)	runner
running	S3 (ing dropped)	runn
quickly	none apply	quickly
easily	none apply	easily
across	none apply	across
fields	S5 (s dropped)	field

Remark: running → runn rather than the “correct” run is a classic stemming over-/under-stemming artifact – our simplified rule set does not undouble the consonant. This is exactly the kind of imprecision real stemmers like Porter’s are carefully engineered to avoid, and why lemmatization is sometimes preferred when precision on word form matters.

Practice Exercise 1.3

Sentence: “The boxes were packed and shipped quickly by the studies team.” Using stoplist {the, were, and, by} and rules S1–S5 above:

- Tokenize and case-fold.
- Remove stop words.
- Apply the stemming rules to each remaining token.

Solution. (a) the, boxes, were, packed, and, shipped, quickly, by, the, studies, team

(b) Removing {the, were, and, by}: boxes, packed, shipped, quickly, studies, team

(c)

Token	Rule applied	Stem
boxes	S5 (s dropped)	boxe
packed	S4 (ed dropped)	pack
shipped	S4 (ed dropped)	shipp
quickly	none apply	quickly
studies	S1 (ies→y)	study
team	none apply	team

Remark: boxes→boxe and shipped→shipp again show why real stemmers add extra rules (e.g., un-doubling a doubled final consonant, restoring a dropped e); our 5-rule set is intentionally simplified for hand practice.

Week 2

Dictionaries, Tolerant Retrieval, and Index Construction

5 Dictionary Data Structures

The dictionary (vocabulary) of an inverted index must support fast lookup of a query term's postings-list pointer. Two common structures:

- **Sorted array / B-tree:** enables binary search, $O(\log |V|)$ lookup, and – crucially – supports *range* and *prefix* queries (e.g., all terms between **hot** and **hotel**), which a hash table cannot.
- **Hash table:** $O(1)$ average lookup for a single exact term, but no support for prefix/wildcard/range queries and it must be resized as the vocabulary grows.

6 Wildcard Queries and the Permuterm Index

Definition 2.1 – Permuterm Index

For each dictionary term t , append a special end-of-word symbol $\$$ and form all *rotations* of the string $t\$$. Each rotation is stored in the dictionary pointing back to t . A wildcard query of the general form $X*Y$ (text X before the wildcard, text Y after) is answered by rotating the query itself to $Y\$X$ and doing a *prefix* search for $Y\$X$ in the rotated dictionary.

Worked Example 2.1 – Permuterm Rotation and Lookup

Take the single dictionary term **money**. Form **money\$** and rotate it one character at a time:

money\$, oney\$m, ney\$mo, ey\$mon, y\$mone, \$money

All six rotations are stored in the permuterm dictionary, each pointing to the original term **money**.

Now answer the wildcard query **mon*y** (matches any word starting with **mon** and ending in **y**). Rewrite $X*Y$ with $X = \text{mon}$, $Y = \text{y}$ as the rotated pattern $Y\$X = \text{y\$mon}$, and search the rotated dictionary for entries with this exact *prefix*. Rotation #5, **y\$mone**, begins with **y\$mon** – match! It points back to **money**, which is returned.

Practice Exercise 2.1

The dictionary contains only the term **hello**.

- Write out all rotations of **hello\$**.
- Which rotation prefix would you search to answer the wildcard query **hel*o**? Show the rewritten pattern and identify the matching rotation.

Solution. (a) **hello\$, ello\$h, llo\$he, lo\$hel, o\$hell, \$hello**.

(b) Query **hel*o** has $X = \text{hel}$, $Y = \text{o}$, so we search for prefix $Y\$X = \text{o$hel}$. Rotation #5, **o\$hell**, begins with **o\$hel** – match, pointing back to **hello**.

Remark: k -gram indexes are a lighter-weight alternative to permuterm indexes: every dictionary term is broken into overlapping k -character substrings (padded with \$ at the boundaries), and each k -gram points to the list of terms containing it. A wildcard query is broken into its own k -grams, and candidate terms are those whose k -gram sets have high overlap with the query's – which is exactly the Jaccard-coefficient idea used below for spelling correction.

7 Spelling Correction

7.1 Edit Distance

Definition 2.2 – Levenshtein (Edit) Distance

The **edit distance** between two strings is the minimum number of single-character insertions, deletions, and substitutions required to transform one string into the other.

Key Formula 2.1 – Edit Distance Dynamic Program

Let s_1 (length n) and s_2 (length m) be the two strings. Define $D[i, j]$ = edit distance between the first i characters of s_1 and the first j characters of s_2 . Then:

$$D[i, 0] = i, \quad D[0, j] = j,$$

$$D[i, j] = \min \begin{cases} D[i-1, j] + 1 & \text{(deletion)} \\ D[i, j-1] + 1 & \text{(insertion)} \\ D[i-1, j-1] + \mathbb{I}[s_1[i] \neq s_2[j]] & \text{(substitution / match)} \end{cases}$$

The answer is $D[n, m]$.

Worked Example 2.2 – Edit Distance by Dynamic Programming Table

Compute the edit distance between **kitten** (s_1) and **sitting** (s_2).

	#	s	i	t	t	i	n	g
#	0	1	2	3	4	5	6	7
k	1	1	2	3	4	5	6	7
i	2	2	1	2	3	4	5	6
t	3	3	2	1	2	3	4	5
t	4	4	3	2	1	2	3	4
e	5	5	4	3	2	2	3	4
n	6	6	5	4	3	3	2	3

Reading off $D[6, 7] = 3$: three edits are needed (substitute **k**→**s**, substitute **e**→**i**, insert **g**), matching **kitten** → **sitten** → **sittin** → **sitting**.

How the last entry was filled, as a template: $D[6, 7] = \min(D[5, 7] + 1, D[6, 6] + 1, D[5, 6] + \mathbb{I}[n \neq g]) = \min(4 + 1, 2 + 1, 3 + 1) = 3$ (the minimum comes from the diagonal move $D[5, 6] + 1$, since $n \neq g$).

Practice Exercise 2.2

Compute the edit distance between **flaw** and **lawn** by filling in the full dynamic-programming table (rows = **flaw**, columns = **lawn**). State the minimum-cost edit sequence.

Solution.

	#	l	a	w	n
#	0	1	2	3	4
f	1	1	2	3	4
l	2	1	2	3	4
a	3	2	1	2	3
w	4	3	2	1	2

$D[4, 4] = 2$. Edit sequence: delete leading f (flaw $\rightarrow$ law), then insert trailing n (law $\rightarrow$ lawn) – 2 edits.

7.2 Jaccard Coefficient for k -gram Candidate Generation

Key Formula 2.2 – Jaccard Coefficient

For two sets X and Y (here, bigram sets of two strings):

$$J(X, Y) = \frac{|X \cap Y|}{|X \cup Y|}, \quad 0 \leq J(X, Y) \leq 1.$$

A spelling corrector uses J as a cheap first filter: only dictionary terms whose k -gram set is sufficiently similar to the misspelled query's k -gram set are passed on to the (more expensive) edit-distance computation.

Worked Example 2.3 – Bigram Jaccard by Hand

Compare **stat** and **stet** using character bigrams (no boundary padding).

Bigrams of **stat**: st, ta, at $\rightarrow$ set {st, ta, at}, size 3.

Bigrams of **stet**: st, te, et $\rightarrow$ set {st, te, et}, size 3.

Intersection: {st}, size 1. Union: {st, ta, at, te, et}, size 5.

$$J(\text{stat}, \text{stet}) = \frac{1}{5} = 0.2$$

Practice Exercise 2.3

Compute the bigram Jaccard coefficient between **night** and **nigth** (a common transposition typo).

Solution. Bigrams of **night**: ni, ig, gh, ht – set size 4.

Bigrams of **nigth**: ni, ig, gt, th – set size 4.

Intersection: {ni, ig}, size 2. Union: {ni, ig, gh, ht, gt, th}, size 6.

$$J = \frac{2}{6} = \frac{1}{3} \approx 0.333$$

8 Index Construction

8.1 Blocked Sort-Based Indexing (BSBI)

When the collection is too large to sort all term–docID pairs in memory at once, BSBI processes the collection in blocks:

Algorithm 2 BSBI-INDEXING

```

1: while documents remain do
2:   read a block of documents that fits in memory
3:   parse it into a list of (term, docID) pairs
4:   sort the pairs by term, then by docID
5:   write the sorted block to disk as a mini inverted index
6: end while
7: merge all the on-disk blocks (multi-way merge) into one final inverted index

```

8.2 Single-Pass In-Memory Indexing (SPIMI)

SPIMI avoids ever sorting explicitly in memory: for each block, it builds postings lists incrementally in a hash table (dictionary term $\rightarrow$ growing postings list) as it streams through the tokens, then dumps the block's dictionary to disk sorted by term, and starts a fresh dictionary for the next block. This is simpler and typically faster than BSBI because there is no separate sort step over raw pairs.

Worked Example 2.4 – Constructing an Inverted Index by Hand

Mini corpus (4 documents; stop words intentionally retained here to show the raw mechanics):

Doc	Text
1	the cat sat
2	the dog sat on the mat
3	cat and dog play
4	the mat is red

Step 1 – emit (term, docID) pairs in document order:

(the,1) (cat,1) (sat,1) (the,2) (dog,2) (sat,2) (on,2) (the,2) (mat,2) (cat,3) (and,3) (dog,3) (play,3)
(the,4) (mat,4) (is,4) (red,4)

Step 2 – sort by term, then by docID:

(and,3) (cat,1) (cat,3) (dog,2) (dog,3) (is,4) (mat,2) (mat,4) (on,2) (play,3) (red,4) (sat,1) (sat,2)
(the,1) (the,2) (the,2) (the,4)

Step 3 – group consecutive same-term entries into a postings list, collapsing repeated docIDs into one entry with a term frequency count:

Term	Postings (docID:freq)
and	3:1
cat	1:1, 3:1
dog	2:1, 3:1
is	4:1
mat	2:1, 4:1
on	2:1
play	3:1
red	4:1
sat	1:1, 2:1
the	1:1, 2:2, 4:1

Notice **the** occurs *twice* within document 2 (“*the* dog sat on *the* mat”), so the two (the,2) pairs collapse into a single postings entry 2:2 recording term frequency 2, rather than two separate entries.

Practice Exercise 2.4

Build the inverted index (postings with term-frequency counts) for this 3-document corpus by hand, following the same three steps:

Doc	Text
1	fox jumps over the fox
2	the dog sleeps
3	fox and dog run

Solution. Term-docID pairs: (fox,1)(jumps,1)(over,1)(the,1)(fox,1) → (the,2)(dog,2)(sleeps,2) → (fox,3)(and,3)(dog,3)(run,3).

Sorted: (and,3)(dog,2)(dog,3)(fox,1)(fox,1)(fox,3)(jumps,1)(over,1)(run,3)(sleeps,2)(the,1)(the,2).

Grouped postings:

Term	Postings (docID:freq)
and	3:1
dog	2:1, 3:1
fox	1:2, 3:1
jumps	1:1
over	1:1
run	3:1
sleeps	2:1
the	1:1, 2:1

(fox appears twice in document 1, collapsing to 1:2.)

8.3 Distributed and Dynamic Indexing (Conceptual Overview)

- **Distributed indexing (MapReduce):** *parsers* (mappers) split the collection into (term, docID) pairs; the pair space is partitioned by term range across machines; *inverters* (reducers) collect all pairs for their assigned term range and write sorted postings – BSBI’s block-and-merge idea, parallelized across a cluster.
- **Dynamic indexing:** new documents cannot cheaply be inserted into one giant sorted index. The standard solution keeps a small, easily-rebuilt **auxiliary index** in memory for recent documents, answers queries against *both* the main and auxiliary index, and periodically **merges** the auxiliary index into the main index (logarithmic merging schedules multiple auxiliary indexes of geometrically increasing size to bound the total re-merging cost).

Week 3

Term Weighting, the Vector Space Model, and Efficient Scoring

9 Parametric and Zone Indexes

Real documents have structure beyond a flat bag of words: a **title**, an **author** field, a **date**, an **abstract**. A **parametric index** lets a query restrict on such fields directly (e.g., **date** = 2023). A **zone** is a named region of free text within a document (title, body, anchor text); a **zone index** keeps a separate postings list per zone so that a match in the title can be treated differently from a match in the body – the basis of *weighted zone scoring* in Section 12.

10 Term Frequency and Document Frequency

Definition 3.1 – tf and df

The **term frequency** $tf_{t,d}$ is the raw count of term t in document d . The **document frequency** df_t is the number of documents in the collection containing t at least once (not to be confused with the *collection frequency*, the total count of t over all documents).

Raw tf is a poor weight on its own: a document with $tf = 20$ is not “20 times as relevant” as one with $tf = 1$. IR systems therefore commonly **dampen** tf , e.g. with a log transform $wf_{t,d} = 1 + \log(tf_{t,d})$ for $tf_{t,d} > 0$ (and 0 otherwise).

10.1 Inverse Document Frequency

Key Formula 3.1 – Inverse Document Frequency

$$idf_t = \log\left(\frac{N}{df_t}\right)$$

where N is the total number of documents in the collection. A term in *every* document has $df_t = N$, so $idf_t = \log(1) = 0$ – it is given *no* weight at all, formalizing the intuition behind stop-word removal (Definition 1.5).

Remark: Any logarithm base may be used (2, 10, e); changing the base multiplies every idf_t (and hence every tf -idf weight) by the same constant $\log_b(a)$, so relative rankings by cosine similarity are completely unaffected. In the worked examples below we deliberately use $\log_2$ with a collection size that is a power of 2, purely so the arithmetic can be done without a calculator.

10.2 tf-idf Weighting

Key Formula 3.2 – tf-idf Weight

$$tf\text{-}idf_{t,d} = tf_{t,d} \times idf_t$$

This single number is high when t occurs often in d (informative about d) and t is rare across the collection (informative in general) – and low whenever either factor is low.

Worked Example 3.1 – Computing idf and tf-idf with Clean Numbers

Collection size $N = 8$. Vocabulary of interest: the, internet, protocol, network, with document frequencies $df(\text{the}) = 8$, $df(\text{internet}) = 4$, $df(\text{protocol}) = 2$, $df(\text{network}) = 1$.

$$idf(\text{the}) = \log_2 \frac{8}{8} = \log_2 1 = 0, \quad idf(\text{internet}) = \log_2 \frac{8}{4} = \log_2 2 = 1,$$

$$idf(\text{protocol}) = \log_2 \frac{8}{2} = \log_2 4 = 2, \quad idf(\text{network}) = \log_2 \frac{8}{1} = \log_2 8 = 3.$$

Document D_3 has raw counts $tf(\text{the}) = 10$, $tf(\text{internet}) = 4$, $tf(\text{protocol}) = 2$, $tf(\text{network}) = 1$. Then:

$$tf\text{-}idf(\text{the}, D_3) = 10 \times 0 = 0, \quad tf\text{-}idf(\text{internet}, D_3) = 4 \times 1 = 4,$$

$$tf\text{-}idf(\text{protocol}, D_3) = 2 \times 2 = 4, \quad tf\text{-}idf(\text{network}, D_3) = 1 \times 3 = 3.$$

So D_3 's tf-idf vector, in the order (the, internet, protocol, network), is $(0, 4, 4, 3)$.

Practice Exercise 3.1

Using the same idf values as Worked Example 3.1, document D_2 has raw counts $tf(\text{the}) = 8$, $tf(\text{internet}) = 2$, $tf(\text{protocol}) = 0$, $tf(\text{network}) = 2$. Compute D_2 's full tf-idf vector.

Solution. $tf\text{-}idf(\text{the}, D_2) = 8 \times 0 = 0$; $tf\text{-}idf(\text{internet}, D_2) = 2 \times 1 = 2$; $tf\text{-}idf(\text{protocol}, D_2) = 0 \times 2 = 0$; $tf\text{-}idf(\text{network}, D_2) = 2 \times 3 = 6$.
 $D_2 = (0, 2, 0, 6)$.

11 The Vector Space Model

Definition 3.2 – Vector Space Model (VSM)

Each document d (and the query q) is represented as a vector in $\mathbb{R}^{|V|}$, one dimension per vocabulary term, with the tf-idf weight as the coordinate. Similarity between q and d is measured as the **cosine** of the angle between their vectors – *not* their Euclidean distance, because cosine similarity is insensitive to document length (a long document that repeats the same topic should not automatically win just because its raw vector is longer).

Key Formula 3.3 – Cosine Similarity

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{\|\vec{q}\| \|\vec{d}\|} = \frac{\sum_t q_t d_t}{\sqrt{\sum_t q_t^2} \sqrt{\sum_t d_t^2}}$$

Equivalently: normalize $\vec{q}$ and $\vec{d}$ to unit vectors first, then take their dot product.

Worked Example 3.2 – Cosine Similarity Intuition in 2 Dimensions

Let $D_1 = (4, 0)$, $D_2 = (0, 3)$, query $Q = (2, 1)$.

$$\|D_1\| = 4, \quad \|D_2\| = 3, \quad \|Q\| = \sqrt{2^2 + 1^2} = \sqrt{5} \approx 2.236$$

$$\cos(Q, D_1) = \frac{2(4) + 1(0)}{4 \times 2.236} = \frac{8}{8.944} \approx 0.894$$

$$\cos(Q, D_2) = \frac{2(0) + 1(3)}{3 \times 2.236} = \frac{3}{6.708} \approx 0.447$$

D_1 ranks above D_2 : its direction is closer to Q 's, even though we never computed a Euclidean distance.

Worked Example 3.3 – Full tf-idf + Cosine Ranking

Reuse $idf(\text{the}) = 0$, $idf(\text{internet}) = 1$, $idf(\text{protocol}) = 2$, $idf(\text{network}) = 3$ and $D_2 = (0, 2, 0, 6)$, $D_3 = (0, 4, 4, 3)$ from above, and add D_1 with raw counts $tf(\text{the}) = 5$, $tf(\text{protocol}) = 3$ (internet, network absent), giving $D_1 = (5 \times 0, 0, 3 \times 2, 0) = (0, 0, 6, 0)$.

Query: “protocol network” – raw query term count 1 each, weighted by idf : $Q = (0, 0, 1 \times 2, 1 \times 3) = (0, 0, 2, 3)$.

Norms:

$$\|D_1\| = \sqrt{6^2} = 6, \quad \|D_2\| = \sqrt{2^2 + 6^2} = \sqrt{40} \approx 6.325, \quad \|D_3\| = \sqrt{4^2 + 4^2 + 3^2} = \sqrt{41} \approx 6.403,$$

$$\|Q\| = \sqrt{2^2 + 3^2} = \sqrt{13} \approx 3.606.$$

Dot products:

$$Q \cdot D_1 = 2(6) + 3(0) = 12, \quad Q \cdot D_2 = 2(0) + 3(6) = 18, \quad Q \cdot D_3 = 2(4) + 3(3) = 17.$$

Cosine similarities:

$$\cos(Q, D_1) = \frac{12}{6 \times 3.606} = \frac{12}{21.63} \approx 0.555,$$

$$\cos(Q, D_2) = \frac{18}{6.325 \times 3.606} = \frac{18}{22.81} \approx 0.789,$$

$$\cos(Q, D_3) = \frac{17}{6.403 \times 3.606} = \frac{17}{23.09} \approx 0.736.$$

Ranking: D_2 (0.789) > D_3 (0.736) > D_1 (0.555).

*Remark: D_2 ranks above D_3 even though D_2 has zero weight on **protocol** while D_3 matches both query terms! Cosine similarity depends on the overall direction of the document vector, not simply on how many query terms it contains – D_2 's very high weight on **network** ($idf = 3$) dominates the dot product and gives it a slightly larger numerator relative to its norm. This is a genuine, well-known property of the vector space model, not an error.*

Practice Exercise 3.2

Using the same three document vectors $D_1 = (0, 0, 6, 0)$, $D_2 = (0, 2, 0, 6)$, $D_3 = (0, 4, 4, 3)$ (order: the, internet, protocol, network), re-rank them for the query “internet protocol”, i.e. $Q = (0, 1 \times 1, 1 \times 2, 0) = (0, 1, 2, 0)$. Compute all three cosine similarities and give the final ranking.

Solution. $\|Q\| = \sqrt{1^2 + 2^2} = \sqrt{5} \approx 2.236.$

$$Q \cdot D_1 = 1(0) + 2(6) = 12; \quad \cos(Q, D_1) = \frac{12}{6 \times 2.236} = \frac{12}{13.42} \approx 0.894.$$

$$Q \cdot D_2 = 1(2) + 2(0) = 2; \quad \cos(Q, D_2) = \frac{2}{6.325 \times 2.236} = \frac{2}{14.14} \approx 0.141.$$

$$Q \cdot D_3 = 1(4) + 2(4) = 12; \quad \cos(Q, D_3) = \frac{12}{6.403 \times 2.236} = \frac{12}{14.32} \approx 0.838.$$

Ranking: D_1 (0.894) > D_3 (0.838) > D_2 (0.141). (Here D_2 drops to last, since it has no weight on protocol and only a small internet weight relative to its large network-dominated norm.)

12 Efficient Scoring

Computing full cosine similarity against every document in the collection is wasteful when only a handful of documents can possibly score well. Real systems use several tricks:

- **Document-at-a-time scoring:** iterate through postings lists in parallel by increasing docID, accumulating each document’s full score before moving to the next docID (requires all postings sorted by docID – exactly the structure used for Boolean merging in Week 1).
- **Term-at-a-time scoring:** process one query term’s entire postings list at a time, accumulating partial scores per document in an accumulator array, then move to the next term. Simpler to reason about, but needs an accumulator per candidate document.
- **Tiered indexes / champion lists:** keep a short list of the highest-*tf-idf* documents for each term (a “champion list”); try to satisfy the query from these small lists first, only falling back to the full postings lists if too few results are found.

12.1 Weighted Zone Scoring

Key Formula 3.4 – Weighted Zone Scoring

Given ℓ zones with weights $g_1, \dots, g_\ell$ (typically $\sum_i g_i = 1$) and a Boolean match indicator $s_i \in \{0, 1\}$ for each zone (does the query term/expression match in zone i ?):

$$\text{score} = \sum_{i=1}^{\ell} g_i s_i$$

Practice Exercise 3.3

A search engine uses two zones, title (weight $g_{\text{title}} = 0.6$) and body (weight $g_{\text{body}} = 0.4$), for the query term **java**. Three documents match as follows: D_1 (title only), D_2 (body only), D_3 (both title and body). Rank D_1, D_2, D_3 by weighted zone score.

Solution. $\text{score}(D_1) = 0.6(1) + 0.4(0) = 0.6$; $\text{score}(D_2) = 0.6(0) + 0.4(1) = 0.4$; $\text{score}(D_3) = 0.6(1) + 0.4(1) = 1.0$.

Ranking: D_3 (1.0) > D_1 (0.6) > D_2 (0.4).

12.2 The BM25 Scoring Function

BM25 (“Best Match 25”) is the classic probabilistically-motivated scoring function used as the industry-standard baseline ranking function, refining raw *tf-idf* with document-length normalization.

Key Formula 3.5 – Okapi BM25

$$\text{score}(D, Q) = \sum_{t \in Q} \text{idf}(t) \cdot \frac{f(t, D) (k_1 + 1)}{f(t, D) + k_1 \left(1 - b + b \frac{|D|}{\text{avgdl}} \right)}$$

where $f(t, D)$ is the raw term frequency of t in D , $|D|$ is the length of D , avgdl is the average document length in the collection, and $\text{idf}(t) = \ln \frac{N - df_t + 0.5}{df_t + 0.5}$. The free parameters k_1 (typically ≈ 1.2) control term-frequency saturation, and b (typically 0.75) controls the strength of length normalization.

Remark: As $f(t, D) \rightarrow \infty$, the term's contribution saturates at $\text{idf}(t) \cdot (k_1 + 1)$ rather than growing without bound as raw $tf\text{-idf}$ would – this is the key mathematical improvement BM25 makes over Key Formula 3.2.

Worked Example 3.4 – Computing a BM25 Score

Collection: $N = 10$, avgdl = 50, $k_1 = 1.2$, $b = 0.75$. Query terms network ($df = 2$) and protocol ($df = 3$). Document D has $|D| = 60$, $f(\text{network}, D) = 3$, $f(\text{protocol}, D) = 1$.

Step 1 – idf:

$$\text{idf}(\text{network}) = \ln \frac{10 - 2 + 0.5}{2 + 0.5} = \ln \frac{8.5}{2.5} = \ln 3.4 \approx 1.2238$$

$$\text{idf}(\text{protocol}) = \ln \frac{10 - 3 + 0.5}{3 + 0.5} = \ln \frac{7.5}{3.5} = \ln 2.1429 \approx 0.7621$$

Step 2 – length-normalization factor (same for every term in this document):

$$1 - b + b \frac{|D|}{\text{avgdl}} = 1 - 0.75 + 0.75 \times \frac{60}{50} = 0.25 + 0.9 = 1.15, \quad k_1 \times 1.15 = 1.38$$

Step 3 – per-term score:

$$\text{network} : \frac{3(1.2 + 1)}{3 + 1.38} = \frac{6.6}{4.38} \approx 1.5068 \Rightarrow 1.2238 \times 1.5068 \approx 1.8441$$

$$\text{protocol} : \frac{1(2.2)}{1 + 1.38} = \frac{2.2}{2.38} \approx 0.9244 \Rightarrow 0.7621 \times 0.9244 \approx 0.7045$$

Step 4 – total:

$$\text{score}(D, Q) = 1.8441 + 0.7045 \approx 2.5486$$

Practice Exercise 3.4

Using the same collection statistics and idf values as Worked Example 3.4 ($\text{idf}(\text{network}) \approx 1.2238$, $\text{idf}(\text{protocol}) \approx 0.7621$, $k_1 = 1.2$, $b = 0.75$, avgdl = 50), a different document D' has $|D'| = 40$, $f(\text{network}, D') = 1$, $f(\text{protocol}, D') = 2$. Compute $\text{score}(D', Q)$.

Solution. Length factor: $1 - 0.75 + 0.75 \times \frac{40}{50} = 0.25 + 0.6 = 0.85$; $k_1 \times 0.85 = 1.02$.

$$\text{network } (f = 1): \frac{1(2.2)}{1 + 1.02} = \frac{2.2}{2.02} \approx 1.0891 \Rightarrow 1.2238 \times 1.0891 \approx 1.3330.$$

$$\text{protocol } (f = 2): \frac{2(2.2)}{2 + 1.02} = \frac{4.4}{3.02} \approx 1.4570 \Rightarrow 0.7621 \times 1.4570 \approx 1.1103.$$

Total: $\text{score}(D', Q) \approx 1.3330 + 1.1103 = 2.4433$.

Week 4

Evaluation, Relevance Feedback, and Query Expansion

13 Evaluating an Unranked Retrieval Set

Definition 4.1 – Precision and Recall

Given a query, a system's returned set, and a ground-truth relevance judgment, define: TP = relevant & retrieved, FP = non-relevant & retrieved, FN = relevant & not retrieved.

$$\text{Precision } P = \frac{TP}{TP + FP}, \quad \text{Recall } R = \frac{TP}{TP + FN}$$

Key Formula 4.1 – F_1 Measure

$$F_1 = \frac{2PR}{P + R}$$

the harmonic mean of precision and recall – chosen over the arithmetic mean because it penalizes a system that scores very low on *either* P or R much more heavily.

Worked Example 4.1 – Precision, Recall, F_1 from a Contingency Table

A collection has 10 documents relevant to a query. A system retrieves 8 documents, of which 5 are relevant.

$$\begin{aligned} TP &= 5, \quad FP = 8 - 5 = 3, \quad FN = 10 - 5 = 5 \\ P &= \frac{5}{5 + 3} = \frac{5}{8} = 0.625, \quad R = \frac{5}{5 + 5} = \frac{5}{10} = 0.5 \\ F_1 &= \frac{2(0.625)(0.5)}{0.625 + 0.5} = \frac{0.625}{1.125} \approx 0.5556 \end{aligned}$$

Practice Exercise 4.1

A collection has 12 documents relevant to a query. A system retrieves 9 documents, of which 6 are relevant. Compute P , R , and F_1 .

Solution. $TP = 6$, $FP = 9 - 6 = 3$, $FN = 12 - 6 = 6$. $P = \frac{6}{9} \approx 0.667$, $R = \frac{6}{12} = 0.5$.
 $F_1 = \frac{2(0.667)(0.5)}{0.667 + 0.5} = \frac{0.667}{1.167} \approx 0.5715$.

14 Evaluating a Ranked Retrieval List

14.1 Average Precision and Mean Average Precision

Key Formula 4.2 – Average Precision (AP)

For one query with R total relevant documents in the collection, let $P(k)$ be the precision computed at the rank k of each successive relevant document retrieved (documents never retrieved contribute 0):

$$AP = \frac{1}{R} \sum_{\text{relevant } d} P(\text{rank of } d)$$

Mean Average Precision (MAP) is the mean of AP over a set of queries:

$$MAP = \frac{1}{|Q|} \sum_{q \in Q} AP(q)$$

Worked Example 4.2 – Computing AP and MAP

Query 1: 6 total relevant documents exist. The ranked results (R = relevant, N = not) are:

$$\underbrace{R}_1 \underbrace{N}_2 \underbrace{R}_3 \underbrace{R}_4 \underbrace{N}_5 \underbrace{N}_6 \underbrace{R}_7 \underbrace{N}_8 \underbrace{N}_9 \underbrace{R}_{10}$$

Precision at each relevant rank: $P(1) = \frac{1}{1} = 1.0$, $P(3) = \frac{2}{3} \approx 0.667$, $P(4) = \frac{3}{4} = 0.75$, $P(7) = \frac{4}{7} \approx 0.571$, $P(10) = \frac{5}{10} = 0.5$. Only 5 of the 6 relevant documents were retrieved at all, so the 6th contributes 0.

$$AP_1 = \frac{1.0 + 0.667 + 0.75 + 0.571 + 0.5 + 0}{6} = \frac{3.488}{6} \approx 0.5814$$

Query 2: 2 total relevant documents, ranked results: $N R R N N$. $P(2) = \frac{1}{2} = 0.5$, $P(3) = \frac{2}{3} \approx 0.667$.

$$AP_2 = \frac{0.5 + 0.667}{2} = \frac{1.167}{2} \approx 0.5835$$

$$MAP = \frac{AP_1 + AP_2}{2} = \frac{0.5814 + 0.5835}{2} \approx 0.5824$$

Practice Exercise 4.2

A query has 4 total relevant documents. The ranked results are: $R R N N R N$ (all 3 relevant documents that exist among these six ranks were retrieved, but a 4th relevant document elsewhere in the collection was never retrieved). Compute AP for this query.

Solution. $P(1) = \frac{1}{1} = 1.0$, $P(2) = \frac{2}{2} = 1.0$, $P(5) = \frac{3}{5} = 0.6$; the 4th relevant document contributes 0.

$$AP = \frac{1.0 + 1.0 + 0.6 + 0}{4} = \frac{2.6}{4} = 0.65$$

14.2 Normalized Discounted Cumulative Gain (nDCG)

Key Formula 4.3 – DCG and nDCG

With graded relevance scores $rel_1, rel_2, \dots, rel_p$ at ranks $1, \dots, p$:

$$DCG_p = rel_1 + \sum_{i=2}^p \frac{rel_i}{\log_2 i}$$

$IDCG_p$ is DCG_p computed on the *ideal* ranking (relevance scores sorted in descending order).

Then:

$$nDCG_p = \frac{DCG_p}{IDCG_p} \in [0, 1]$$

Useful values: $\log_2 2 = 1$, $\log_2 3 \approx 1.585$, $\log_2 4 = 2$, $\log_2 5 \approx 2.322$, $\log_2 6 \approx 2.585$

Worked Example 4.3 – Computing nDCG

Ranked list of 5 documents with graded relevance $rel = [3, 2, 3, 0, 1]$.

Actual DCG:

$$\begin{aligned} DCG_5 &= 3 + \frac{2}{\log_2 2} + \frac{3}{\log_2 3} + \frac{0}{\log_2 4} + \frac{1}{\log_2 5} = 3 + \frac{2}{1} + \frac{3}{1.585} + \frac{0}{2} + \frac{1}{2.322} \\ &= 3 + 2 + 1.893 + 0 + 0.431 = 7.324 \end{aligned}$$

Ideal ranking (sort relevance descending): $[3, 3, 2, 1, 0]$.

$$\begin{aligned} IDCG_5 &= 3 + \frac{3}{1} + \frac{2}{1.585} + \frac{1}{2} + \frac{0}{2.322} = 3 + 3 + 1.262 + 0.5 + 0 = 7.762 \\ nDCG_5 &= \frac{7.324}{7.762} \approx 0.9436 \end{aligned}$$

Practice Exercise 4.3

Ranked list of 6 documents with graded relevance $rel = [2, 3, 2, 0, 1, 2]$. Compute DCG_6 , $IDCG_6$, and $nDCG_6$.

Solution.

$$\begin{aligned} DCG_6 &= 2 + \frac{3}{\log_2 2} + \frac{2}{\log_2 3} + \frac{0}{\log_2 4} + \frac{1}{\log_2 5} + \frac{2}{\log_2 6} = 2 + 3 + \frac{2}{1.585} + 0 + \frac{1}{2.322} + \frac{2}{2.585} \\ &= 2 + 3 + 1.262 + 0 + 0.431 + 0.774 = 7.467 \end{aligned}$$

Ideal order (descending): $[3, 2, 2, 2, 1, 0]$.

$$\begin{aligned} IDCG_6 &= 3 + \frac{2}{1} + \frac{2}{1.585} + \frac{2}{2} + \frac{1}{2.322} + \frac{0}{2.585} = 3 + 2 + 1.262 + 1 + 0.431 + 0 = 7.693 \\ nDCG_6 &= \frac{7.467}{7.693} \approx 0.9706 \end{aligned}$$

15 Relevance Feedback: The Rocchio Algorithm

Definition 4.2 – Relevance Feedback

After an initial query, the user marks some returned documents *relevant* and others *non-relevant*. The system uses these judgments to build an improved query and re-run the search – without the user ever typing a new query.

Key Formula 4.4 – Rocchio’s Formula

$$\vec{q}_m = \alpha \vec{q}_0 + \beta \frac{1}{|D_r|} \sum_{\vec{d} \in D_r} \vec{d} - \gamma \frac{1}{|D_{nr}|} \sum_{\vec{d} \in D_{nr}} \vec{d}$$

where $\vec{q}_0$ is the original query vector, D_r and D_{nr} are the sets of documents (as tf-idf vectors) marked relevant and non-relevant, and $\alpha, \beta, \gamma \geq 0$ trade off trust in the original query versus the feedback (common published defaults: $\alpha = 1$, $\beta = 0.75$, $\gamma = 0.15$; we use $\gamma = 0.25$ below purely to keep the hand-arithmetic clean).

Worked Example 4.4 – Rocchio Update by Hand

$\vec{q}_0 = (1, 0, 1)$. One relevant document $\vec{d}_1 = (1, 1, 0)$, one non-relevant document $\vec{d}_2 = (0, 1, 1)$. Use $\alpha = 1, \beta = 0.75, \gamma = 0.25$.

Since there is only one document in each set, the “mean” is just that document itself:

$$\begin{aligned} \vec{q}_m &= 1 \cdot (1, 0, 1) + 0.75 \cdot (1, 1, 0) - 0.25 \cdot (0, 1, 1) \\ &= (1, 0, 1) + (0.75, 0.75, 0) - (0, 0.25, 0.25) = (1 + 0.75 - 0, 0 + 0.75 - 0.25, 1 + 0 - 0.25) \\ &\vec{q}_m = (1.75, 0.5, 0.75) \end{aligned}$$

Notice the weight on dimension 2 rose from 0 to 0.5 (pulled toward the relevant document) while dimension 3 fell from 1 to 0.75 (pushed away from the non-relevant document, which also had a 1 there).

Practice Exercise 4.4

$\vec{q}_0 = (0, 1, 1)$. Two relevant documents $\vec{d}_1 = (1, 1, 0)$ and $\vec{d}_2 = (1, 0, 1)$; one non-relevant document $\vec{d}_3 = (0, 0, 1)$. Using $\alpha = 1, \beta = 0.75, \gamma = 0.25$, compute $\vec{q}_m$. (Remember D_r ’s contribution uses the *mean* of $\vec{d}_1, \vec{d}_2$.)

Solution. Mean of relevant set: $\frac{(1, 1, 0) + (1, 0, 1)}{2} = (1, 0.5, 0.5)$.

$$\begin{aligned} \vec{q}_m &= 1 \cdot (0, 1, 1) + 0.75 \cdot (1, 0.5, 0.5) - 0.25 \cdot (0, 0, 1) \\ &= (0, 1, 1) + (0.75, 0.375, 0.375) - (0, 0, 0.25) = (0.75, 1.375, 1.125) \end{aligned}$$

16 Probabilistic Relevance Feedback (Preview)

Vector-space (Rocchio) feedback moves the *query vector* toward relevant documents and away from non-relevant ones – a purely geometric operation. **Probabilistic** relevance feedback instead updates the *term-weight statistics* themselves: for each query term t , the counts of how often t appears in the (now-judged) relevant vs. non-relevant documents are refined, and a new probabilistic weight

(Key Formula 5.1, Week 5) is recomputed from these updated counts. The two approaches are compared directly in Section 20 once the probabilistic model has been introduced.

17 Query Expansion

Definition 4.3 – Query Expansion

Query expansion adds new terms to a query, either from a hand-built **thesaurus** (e.g., automatically adding **automobile** when the user types **car**) or from an **automatically derived** thesaurus built by clustering terms that co-occur frequently in the same documents or the same local context windows across the collection.

Query expansion and relevance feedback solve the same underlying problem – the user’s original keywords are an imperfect proxy for their information need – but query expansion modifies the query *before* seeing any results (using global collection statistics), whereas relevance feedback modifies it *after* seeing which specific documents the user judged relevant.

Week 5

Probabilistic Information Retrieval and Language Models for IR

18 The Probability Ranking Principle

Definition 5.1 – Probability Ranking Principle (PRP)

If documents are ranked in decreasing order of their probability of relevance to the query, $P(R = 1 \mid d, q)$, given all the evidence available, then the overall effectiveness of the system to its users is the best obtainable.

The PRP is the theoretical justification for *every* ranking function in this course: tf-idf/cosine, BM25, and language-model scores are all, in different ways, attempts to *approximate* $P(R = 1 \mid d, q)$.

19 The Binary Independence Model (BIM)

Definition 5.2 – Binary Independence Model

Documents and queries are represented as binary term-incidence vectors (as in Week 1’s incidence matrix). “Independence” means terms are assumed to occur *independently of one another* within the relevant documents, and independently within the non-relevant documents (a strong simplifying assumption, but one that leads to a tractable, useful ranking formula).

For a single query term t , define, over the whole judged collection: N = total documents, R = total relevant documents, n_t = documents containing t , r_t = relevant documents containing t . This gives a 2×2 contingency table:

	Relevant	Non-relevant
t present	r_t	$n_t - r_t$
t absent	$R - r_t$	$(N - n_t) - (R - r_t)$

Key Formula 5.1 – RSJ (Robertson–Sparck Jones) Term Weight

Let $p_t = r_t/R$ (probability a relevant document contains t) and $u_t = (n_t - r_t)/(N - R)$ (probability a non-relevant document contains t). The retrieval status value contributed by matching term t is:

$$c_t = \log \frac{p_t(1 - u_t)}{u_t(1 - p_t)}$$

A document’s overall score is the sum of c_t over all query terms t that it contains. Larger c_t means t is a much better discriminator of relevance (much more likely in relevant than non-relevant documents).

Worked Example 5.1 – Computing an RSJ Term Weight

Collection: $N = 1000$ documents, $R = 25$ documents judged relevant to a query. Term t appears in $n_t = 50$ documents total, of which $r_t = 20$ are relevant.

$$p_t = \frac{r_t}{R} = \frac{20}{25} = 0.8, \quad u_t = \frac{n_t - r_t}{N - R} = \frac{50 - 20}{1000 - 25} = \frac{30}{975} \approx 0.03077$$

$$c_t = \log_2 \frac{0.8 \times (1 - 0.03077)}{0.03077 \times (1 - 0.8)} = \log_2 \frac{0.8 \times 0.96923}{0.03077 \times 0.2} = \log_2 \frac{0.77538}{0.006154}$$

$$= \log_2(126.0) \approx 6.977 \text{ bits}$$

A document containing this single term t would receive a document score of ≈ 6.98 .

Practice Exercise 5.1

Collection: $N = 2000$, $R = 40$. Term t appears in $n_t = 80$ documents, of which $r_t = 30$ are relevant. Compute p_t , u_t , and c_t (base 2).

Solution.

$$p_t = \frac{30}{40} = 0.75, \quad u_t = \frac{80 - 30}{2000 - 40} = \frac{50}{1960} \approx 0.02551$$

$$c_t = \log_2 \frac{0.75 \times 0.97449}{0.02551 \times 0.25} = \log_2 \frac{0.73087}{0.006378} = \log_2(114.6) \approx 6.84 \text{ bits}$$

20 Probabilistic Relevance Feedback

Once a user has judged some retrieved documents relevant/non-relevant, the contingency-table counts r_t and n_t for every query term can be *re-estimated* directly from those judgments, and the RSJ weight c_t recomputed with Key Formula 5.1 – this is probabilistic relevance feedback.

Definition 5.3 – Vector Space vs. Probabilistic Relevance Feedback

Rocchio (vector space) feedback, Key Formula 4.4, is a purely *geometric* operation: it nudges the query vector's coordinates toward the centroid of relevant documents and away from non-relevant ones, in the tf-idf vector space. It has no explicit probabilistic interpretation – β and γ are tuned empirically, not derived from a model of relevance.

Probabilistic (BIM) feedback re-estimates the *parameters of an explicit generative model of relevance* (p_t and u_t for every term) directly from the judged documents, using Key Formula 5.1. The weight update is a direct consequence of Bayesian updating, not a geometric heuristic.

	Vector space (Rocchio)	Probabilistic (BIM)
What is updated	Query vector coordinates	Per-term probabilities p_t, u_t
Theoretical grounding	Heuristic / geometric	Probability Ranking Principle
Uses non-relevant docs how	Subtracts their (weighted) vector	Updates u_t (false-positive rate)
New terms discovered via	Any term with nonzero weight in D_r	Any term whose n_t, r_t can be counted

21 Language Models for Information Retrieval

Definition 5.4 – Query Likelihood Model

For each document d , build a **unigram language model** M_d : a probability distribution over vocabulary terms estimated from d 's own text. Rank documents by $P(q | M_d)$ – the probability that d 's language model would have generated the query q . Intuitively: *which document “sounds like” it was written by someone who would also write this query?*

Key Formula 5.2 – MLE Unigram Model and the Zero-Probability Problem

$$P_{\text{mle}}(t | d) = \frac{tf_{t,d}}{|d|}$$

The raw maximum-likelihood estimate assigns probability **exactly zero** to any query term absent from d – a single missing term would make the entire product $P(q | d) = \prod_t P(t | d)$ zero, however well d matches the rest of the query. This is why language models for IR always **smooth** the estimate with a collection-wide background model.

21.1 Jelinek–Mercer Smoothing

Key Formula 5.3 – Jelinek–Mercer (Linear Interpolation) Smoothing

$$P_{JM}(t | d) = (1 - \lambda) P_{\text{mle}}(t | d) + \lambda P(t | C)$$

where $P(t | C)$ is the term's relative frequency in the whole collection C (the background model) and $\lambda \in [0, 1]$ controls the smoothing strength. The query likelihood is $P(q | d) = \prod_{t \in q} P_{JM}(t | d)$.

Worked Example 5.2 – Query Likelihood with Jelinek–Mercer Smoothing

Vocabulary $\{\text{cat}, \text{dog}, \text{mouse}\}$. Document d : “cat cat dog”, so $|d| = 3$, $tf(\text{cat}, d) = 2$, $tf(\text{dog}, d) = 1$, $tf(\text{mouse}, d) = 0$.

Collection background probabilities (estimated from the whole corpus): $P(\text{cat} | C) = 0.2$, $P(\text{dog} | C) = 0.3$, $P(\text{mouse} | C) = 0.1$.

Query $q = \text{“cat mouse”}$. Use $\lambda = 0.4$.

MLE estimates: $P_{\text{mle}}(\text{cat} | d) = \frac{2}{3} \approx 0.667$; $P_{\text{mle}}(\text{mouse} | d) = \frac{0}{3} = 0$.

Smoothed estimates:

$$P_{JM}(\text{cat} | d) = 0.6 \times 0.667 + 0.4 \times 0.2 = 0.400 + 0.080 = 0.480$$

$$P_{JM}(\text{mouse} | d) = 0.6 \times 0 + 0.4 \times 0.1 = 0 + 0.040 = 0.040$$

$$P(q | d) = P_{JM}(\text{cat} | d) \times P_{JM}(\text{mouse} | d) = 0.480 \times 0.040 = 0.0192$$

Note that without smoothing, $P(q | d)$ would have been exactly 0 (because **mouse** never occurs in d) – smoothing is what makes the model usable at all.

21.2 Dirichlet Smoothing

Key Formula 5.4 – Dirichlet-Prior Smoothing

$$P_{\text{Dir}}(t \mid d) = \frac{tf_{t,d} + \mu P(t \mid C)}{|d| + \mu}$$

where μ is a single free parameter (behaving like a pseudo-count of “extra” collection-flavored word occurrences added to every document). Unlike Jelinek–Mercer’s fixed λ , Dirichlet smoothing automatically smooths short documents more heavily than long ones, since μ is a constant while $|d|$ varies.

Practice Exercise 5.2

Using the same document d (cat cat dog, $|d| = 3$) and background probabilities as Worked Example 5.2, apply **Dirichlet smoothing with** $\mu = 5$ instead of Jelinek–Mercer, and recompute $P(q \mid d)$ for $q = \text{“cat mouse”}$.

Solution.

$$P_{\text{Dir}}(\text{cat} \mid d) = \frac{2 + 5(0.2)}{3 + 5} = \frac{2 + 1}{8} = \frac{3}{8} = 0.375$$

$$P_{\text{Dir}}(\text{mouse} \mid d) = \frac{0 + 5(0.1)}{3 + 5} = \frac{0.5}{8} = 0.0625$$

$$P(q \mid d) = 0.375 \times 0.0625 = 0.02344$$

Remark: Comparing the two smoothing methods on the identical document and query: $P_{\text{JM}}(q \mid d) = 0.0192$ vs. $P_{\text{Dir}}(q \mid d) = 0.02344$ – different smoothing methods produce different absolute probabilities, which is expected (they encode different smoothing assumptions); only the ranking across documents needs to be sensible, not the absolute magnitude.

Week 6

Text Classification and Naïve Bayes

22 The Text Classification Problem

Definition 6.1 – Text Classification

Given a fixed set of classes $\mathbb{C} = \{c_1, \dots, c_J\}$ and a **training set** of documents each pre-labeled with a class, learn a classifier that assigns the correct class to a **new, unseen** document. This is *supervised* learning, in contrast to the unsupervised indexing/ranking methods of Weeks 1–5.

Text classification underlies spam filtering, sentiment analysis, language identification, and (as we will see in Week 7) parts of web search quality control.

23 The Naive Bayes Generative Model

Definition 6.2 – Multinomial Naïve Bayes

Naive Bayes models each document as generated by: (1) picking a class c with probability $P(c)$, then (2) generating each word position independently from c 's own word distribution $P(t | c)$. Two independence assumptions make this tractable and give the model its name:

- **Conditional independence:** the presence of one word does not affect the probability of another, given the class.
- **Positional independence:** $P(t | c)$ does not depend on a word's position in the document.

Both assumptions are literally false for real text – yet the resulting classifier is simple, fast, and surprisingly effective.

Key Formula 6.1 – MAP Classification Rule

$$c_{\text{map}} = \arg \max_{c \in \mathbb{C}} P(c) \prod_{t \in d} P(t | c)^{tf_{t,d}}$$

(equivalently, working in log space to avoid numeric underflow: $\arg \max_c [\log P(c) + \sum_{t \in d} tf_{t,d} \log P(t | c)]$).

Key Formula 6.2 – Parameter Estimation with Add-One (Laplace) Smoothing

$$\hat{P}(c) = \frac{N_c}{N}, \quad \hat{P}(t | c) = \frac{T_{ct} + 1}{(\sum_{t' \in V} T_{ct'}) + |V|}$$

where N_c is the number of training documents in class c , N is the total number of training documents, T_{ct} is the number of occurrences of term t in training documents of class c , and $|V|$ is the vocabulary size. Add-one smoothing is essential: without it, any term never seen in class c 's training data would force $P(t | c) = 0$, which would zero out the *entire* product for any test document containing that term – exactly the zero-probability problem of Key Formula 5.2.

Worked Example 6.1 – The Classic China/Japan Naïve Bayes Example

Training set (class c = “China”, class j = “not-China”):

Doc	Text	Class
1	Chinese Beijing Chinese	c
2	Chinese Chinese Shanghai	c
3	Chinese Macao	c
4	Tokyo Japan Chinese	j

Test document 5: “Chinese Chinese Chinese Tokyo Japan” – classify it.

Vocabulary $V = \{\text{Chinese, Beijing, Shanghai, Macao, Tokyo, Japan}\}$, $|V| = 6$.

Priors: $N = 4$ training docs, $N_c = 3$, $N_j = 1$.

$$\hat{P}(c) = \frac{3}{4} = 0.75, \quad \hat{P}(j) = \frac{1}{4} = 0.25$$

Class c token counts (docs 1–3, total length 8): Chinese = 5, Beijing = 1, Shanghai = 1, Macao = 1, Tokyo = 0, Japan = 0.

$$\hat{P}(\text{Chinese} | c) = \frac{5+1}{8+6} = \frac{6}{14} = \frac{3}{7}, \quad \hat{P}(\text{Tokyo} | c) = \hat{P}(\text{Japan} | c) = \frac{0+1}{8+6} = \frac{1}{14}$$

Class j token counts (doc 4, total length 3): Chinese = 1, Tokyo = 1, Japan = 1.

$$\hat{P}(\text{Chinese} | j) = \frac{1+1}{3+6} = \frac{2}{9}, \quad \hat{P}(\text{Tokyo} | j) = \hat{P}(\text{Japan} | j) = \frac{1+1}{3+6} = \frac{2}{9}$$

Score the test document (Chinese×3, Tokyo×1, Japan×1):

$$\begin{aligned} \text{Score}(c) &= \hat{P}(c) \cdot \hat{P}(\text{Chinese} | c)^3 \cdot \hat{P}(\text{Tokyo} | c) \cdot \hat{P}(\text{Japan} | c) = 0.75 \times \left(\frac{3}{7}\right)^3 \times \frac{1}{14} \times \frac{1}{14} \\ &= 0.75 \times 0.07872 \times \frac{1}{196} \approx 0.75 \times 0.07872 \times 0.005102 \approx 3.012 \times 10^{-4} \\ \text{Score}(j) &= \hat{P}(j) \cdot \hat{P}(\text{Chinese} | j)^3 \cdot \hat{P}(\text{Tokyo} | j) \cdot \hat{P}(\text{Japan} | j) = 0.25 \times \left(\frac{2}{9}\right)^5 \\ &= 0.25 \times 0.0005419 \approx 1.355 \times 10^{-4} \end{aligned}$$

Since $\text{Score}(c) \approx 3.01 \times 10^{-4} > \text{Score}(j) \approx 1.36 \times 10^{-4}$, the classifier assigns **class c (China)** to document 5. Normalizing: $P(c | d_5) = \frac{3.012}{3.012 + 1.355} \approx 0.690$, i.e., about 69% confidence.

Practice Exercise 6.1 – Sports vs. Politics

Training set:

Doc	Text	Class
1	goal match win	Sports
2	goal team win match	Sports
3	vote election win	Politics
4	election vote president	Politics

Test document 5: “win vote election”. Using add-one smoothing, compute $\hat{P}(\text{Sports})$, $\hat{P}(\text{Politics})$, the relevant per-class term probabilities, and classify document 5.

Solution. Vocabulary $V = \{\text{goal, match, win, team, vote, election, president}\}$, $|V| = 7$. $N = 4$, $N_{\text{Sports}} = N_{\text{Politics}} = 2$, so $\hat{P}(\text{Sports}) = \hat{P}(\text{Politics}) = 0.5$.

Sports token counts (total length 7): goal = 2, match = 2, win = 2, team = 1.

$$\hat{P}(\text{win} \mid \text{Sports}) = \frac{2+1}{7+7} = \frac{3}{14}, \quad \hat{P}(\text{vote} \mid \text{Sports}) = \hat{P}(\text{election} \mid \text{Sports}) = \frac{0+1}{14} = \frac{1}{14}$$

Politics token counts (total length 6): vote = 2, election = 2, win = 1, president = 1.

$$\hat{P}(\text{win} \mid \text{Politics}) = \frac{1+1}{6+7} = \frac{2}{13}, \quad \hat{P}(\text{vote} \mid \text{Politics}) = \hat{P}(\text{election} \mid \text{Politics}) = \frac{2+1}{13} = \frac{3}{13}$$

Scores for “win vote election”:

$$\text{Score}(\text{Sports}) = 0.5 \times \frac{3}{14} \times \frac{1}{14} \times \frac{1}{14} = 0.5 \times \frac{3}{2744} \approx 0.5 \times 0.001093 \approx 5.47 \times 10^{-4}$$

$$\text{Score}(\text{Politics}) = 0.5 \times \frac{2}{13} \times \frac{3}{13} \times \frac{3}{13} = 0.5 \times \frac{18}{2197} \approx 0.5 \times 0.008193 \approx 4.10 \times 10^{-3}$$

Since $\text{Score}(\text{Politics}) \approx 4.10 \times 10^{-3} \gg \text{Score}(\text{Sports}) \approx 0.547 \times 10^{-3}$ (about $7.5\times$ larger), document 5 is classified as **Politics** – consistent with intuition, since **vote** and **election** are strongly associated with the Politics class.

24 Evaluating a Classifier

Definition 6.3 – Confusion Matrix and Averaging

For each class c , treat classification as a one-vs-rest detection problem and compute $P_c, R_c, F_{1,c}$ exactly as in Definition 4.1 (Week 4), with TP_c = documents correctly assigned to c , FP_c = documents wrongly assigned to c , FN_c = documents belonging to c but assigned elsewhere.

- **Macro-averaging:** compute P_c (or F_1) for every class separately, then take the simple average – every class counts equally, regardless of size.
- **Micro-averaging:** pool all TP , FP , FN counts across all classes first, then compute a single P, R, F_1 – larger classes dominate the result.

Week 7

Web Search Basics and Link Analysis

25 Characteristics of the Web as a Retrieval Problem

The web differs from a classical, closed IR collection in ways that matter for every model studied so far:

- **Scale and dynamism:** billions of pages, created, edited, and deleted continuously – no fixed N .
- **No quality control:** unlike a curated corpus, anyone can publish anything, motivating spam detection and quality/authority signals (Section 27).
- **Hyperlinks:** a uniquely web-specific source of evidence about a page's importance and topic, unavailable in Weeks 1–6's text-only models.

26 Web Crawling (Conceptual Overview)

Definition 7.1 – Web Crawler

A **web crawler** (spider) starts from a set of **seed URLs**, maintains a **URL frontier** (queue of URLs to fetch), downloads each page, extracts its outgoing links, and adds newly discovered URLs back into the frontier – a graph traversal (conceptually breadth-first) over the web graph. Real crawlers must additionally respect **politeness** (rate-limiting per host, honoring `robots.txt`), prioritize by estimated page importance/freshness, and detect near-duplicate content.

26.1 Estimating Relative Search-Engine Coverage from Overlap

A classic problem (and this week's discussion question) asks how big one search engine's index is *relative to* another's, using only overlap statistics – no crawling required. This is a real historical technique (Bharat & Broder, 1998, used exactly this kind of random-query-sampling overlap analysis to estimate the relative and absolute sizes of early web search engines).

Worked Example 7.1 – Relative Index Size from Overlap

Search engines A and B each sample a large number of pages uniformly at random from their own index. It is observed that 30% of A 's sampled pages are also present in B 's index, while 50% of B 's sampled pages are also present in A 's index. What is $|A|$ relative to $|B|$?

Reasoning: let O be the number of pages in the overlap of the two indexes (pages both engines have indexed). Since 30% of A 's pages fall in this overlap, $O = 0.30 |A|$. Since 50% of B 's pages fall in the *same* overlap, $O = 0.50 |B|$ too. These are two expressions for the same quantity, so:

$$0.30 |A| = 0.50 |B| \quad \implies \quad \frac{|A|}{|B|} = \frac{0.50}{0.30} = \frac{5}{3} \approx 1.667$$

A 's index is about $\frac{5}{3}$ (roughly $1.67\times$) the size of B 's – i.e., $|A| : |B| = 5 : 3$. Intuitively: since a much smaller *fraction* of A (30%) is needed to cover the same absolute overlap that makes up half of B , A must be the larger index.

Practice Exercise 7.1

Search engines C and D sample their indexes at random. 40% of C 's pages are found in D 's index; 20% of D 's pages are found in C 's index. Find $|C| : |D|$, and state which engine has the larger index.

Solution. $O = 0.40 |C| = 0.20 |D| \Rightarrow \frac{|C|}{|D|} = \frac{0.20}{0.40} = 0.5$, i.e., $|C| : |D| = 1 : 2$. Engine D 's index is twice as large as C 's.

27 Link Analysis: PageRank

PageRank is the canonical example of CLO 4's “algebra and probability” methods applied to web search: it can be derived either as the **stationary distribution of a Markov chain** (probability) or as the **eigenvector of a transition matrix** / solution of a linear system (algebra) – and, as this section shows, both routes converge to the identical numbers.

Definition 7.2 – The Random Surfer Model

Imagine a “random surfer” who, at each step, follows a uniformly random outgoing link from the current page with probability d (the **damping factor**), or jumps to a uniformly random page in the whole collection with probability $1 - d$ (modeling a user typing a fresh URL, and guaranteeing the walk cannot get trapped forever inside a link-free cluster). $PR(p)$ is defined as the long-run fraction of time the surfer spends at page p .


Key Formula 7.1 – PageRank

$$PR(p) = \frac{1-d}{N} + d \sum_{q \rightarrow p} \frac{PR(q)}{L(q)}$$

where N is the total number of pages, the sum is over all pages q with a link to p , and $L(q)$ is the number of outgoing links from q . Typically $d = 0.85$. This is solved either **iteratively** (power iteration: start from a uniform guess and repeatedly apply the formula until it stabilizes) or **exactly** (as a system of N linear equations in N unknowns, using $\sum_p PR(p) = 1$).

Worked Example 7.2 – PageRank by Power Iteration and by Exact Algebra

Four pages $\{A, B, C, D\}$ with hyperlinks $A \rightarrow B$, $A \rightarrow C$, $B \rightarrow C$, $C \rightarrow A$, $D \rightarrow C$ (so $L(A) = 2, L(B) = 1, L(C) = 1, L(D) = 1$; note nobody links to D). Use $d = 0.85$, $N = 4$, so the teleport term is $\frac{1-d}{N} = \frac{0.15}{4} = 0.0375$.



Update equations (from the inlinks of each page):

$$PR(A) = 0.0375 + 0.85 PR(C), \quad PR(B) = 0.0375 + 0.85 \frac{PR(A)}{2}$$

$$PR(C) = 0.0375 + 0.85 \left(\frac{PR(A)}{2} + PR(B) + PR(D) \right), \quad PR(D) = 0.0375$$

(D has no inlinks, so its rank is *always* exactly the teleport value – it never changes across iterations.)

Power iteration, starting from $PR_0(A) = PR_0(B) = PR_0(C) = PR_0(D) = 0.25$:

Iteration	$PR(A)$	$PR(B)$	$PR(C)$	$PR(D)$
0	0.2500	0.2500	0.2500	0.2500
1	0.2500	0.1438	0.5688	0.0375
2	0.5209	0.1438	0.2978	0.0375
3	0.2906	0.2589	0.4130	0.0375

(Sample computation, iteration 2: $PR_2(A) = 0.0375 + 0.85(0.5688) = 0.5209$; $PR_2(C) = 0.0375 + 0.85\left(\frac{0.25}{2} + 0.1438 + 0.0375\right) = 0.0375 + 0.85(0.30625) = 0.2978$.) Each row sums to 1, as required.

Exact algebraic solution. Substituting $PR(D) = 0.0375$ and $PR(B) = 0.0375 + 0.425 PR(A)$ into the $PR(C)$ equation, then into the $PR(A)$ equation, gives a single equation in $PR(A)$:

$$PR(A) = 0.1235625 + 0.6683125 PR(A) \implies PR(A) = \frac{0.1235625}{0.3316875} \approx 0.3725$$

$$PR(B) = 0.0375 + 0.425(0.3725) \approx 0.1958, \quad PR(C) = \frac{PR(A) - 0.0375}{0.85} \approx 0.3941, \quad PR(D) = 0.0375$$

Check: $0.3725 + 0.1958 + 0.3941 + 0.0375 = 0.9999 \approx 1$ (rounding). The power-iteration table above is visibly converging toward exactly this vector – C has the highest rank because it receives inlinks from three different pages (A , B , and D).

Practice Exercise 7.2

Three pages $\{X, Y, Z\}$ with links $X \rightarrow Y$, $Y \rightarrow Z$, $Z \rightarrow X$, $Z \rightarrow Y$ (so $L(X) = 1, L(Y) = 1, L(Z) = 2$). Use $d = 0.85$, $N = 3$.

- Write the three PageRank update equations.
- Starting from $PR_0(X) = PR_0(Y) = PR_0(Z) = \frac{1}{3}$, compute PR_1 and PR_2 for all three pages.
- Set up (you do not need to fully re-derive) the exact linear system and confirm the given answer $PR(X) \approx 0.215$, $PR(Y) \approx 0.397$, $PR(Z) \approx 0.388$ is at least *consistent* with the

direction your iterations are moving.

Solution. (a) Teleport term = $\frac{1 - 0.85}{3} = 0.05$.

$$PR(X) = 0.05 + 0.425 PR(Z), \quad PR(Y) = 0.05 + 0.85 PR(X) + 0.425 PR(Z), \quad PR(Z) = 0.05 + 0.85 PR(Y)$$

(b)

Iteration	$PR(X)$	$PR(Y)$	$PR(Z)$
0	0.3333	0.3333	0.3333
1	0.1917	0.4750	0.3333
2	0.1917	0.3546	0.4538

(e.g., $PR_1(Y) = 0.05 + 0.85(0.3333) + 0.425(0.3333) = 0.05 + 0.2833 + 0.1417 = 0.4750$; $PR_2(Z) = 0.05 + 0.85(0.4750) = 0.4538$.)

(c) After 2 iterations Y is still the largest (as in the target answer) and Z is rising toward its target of 0.388 while X has already nearly settled near 0.19–0.22; continuing the iteration a few more rounds converges to the stated (0.215, 0.397, 0.388), confirming consistency.

27.1 HITS: Hubs and Authorities (Brief)

An alternative link-analysis model computes two scores per page: a **hub score** (points to many good authorities) and an **authority score** (is pointed to by many good hubs), updated by mutual reinforcement: $auth(p) \leftarrow \sum_{q \rightarrow p} hub(q)$, then $hub(p) \leftarrow \sum_{p \rightarrow r} auth(r)$, repeated and renormalized – unlike PageRank, HITS scores are computed *per query* over the local neighborhood of pages relevant to that query, not once for the whole web.

Week 8

Course Review and Synthesis

28 The Big Picture: How the Models Connect

Every model in this course answers the same underlying question – *how relevant is document d to query q ?* – with increasing sophistication:

Model	Core idea	Main limitation it fixes vs. the previous model
Boolean (Wk 1)	Exact set membership via AND/OR/NOT	– (baseline)
tf-idf / VSM (Wk 3)	Geometric similarity, graded relevance	Boolean gives no ranking at all
BM25 (Wk 3)	Saturating tf, document-length normalization	raw tf-idf grows without bound and ignores length
BIM / RSJ (Wk 5)	Explicit probability of relevance, from Bayes' rule	VSM has no probabilistic justification
Language Models (Wk 5)	Smoothed generative model of text itself	BIM's binary (0/1) term representation loses tf information
Naïve Bayes (Wk 6)	Same generative machinery, applied to a <i>labeled class</i> instead of a query	extends the LM idea to supervised classification
PageRank (Wk 7)	Query-independent authority from hyperlink structure	all previous models score <i>content only</i> , ignoring the web graph

29 Master Formula Sheet

Concept	Formula
Merge cost	$O(x + y)$
idf	$idf_t = \log(N/df_t)$
tf-idf	$tf \cdot idf_{t,d} = tf_{t,d} \times idf_t$
Cosine similarity	$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \cdot \vec{d}}{\ \vec{q}\ \ \vec{d}\ }$
BM25	$\sum_{t \in Q} idf(t) \frac{f(t, D)(k_1 + 1)}{f(t, D) + k_1(1 - b + b D /avgdl)}$
Precision / Recall / F_1	$P = \frac{TP}{TP + FP}, R = \frac{TP}{TP + FN}, F_1 = \frac{2PR}{P + R}$
AP / MAP	$AP = \frac{1}{R} \sum_{rel\ d} P(\text{rank of } d); MAP = \text{mean of } AP \text{ over queries}$
DCG / nDCG	$DCG_p = rel_1 + \sum_{i=2}^p \frac{rel_i}{\log_2 i}; nDCG_p = DCG_p / IDCG_p$
Rocchio	$\vec{q}_m = \alpha \vec{q}_0 + \beta \overline{D_r} - \gamma \overline{D_{nr}}$
RSJ weight	$c_t = \log \frac{p_t(1 - u_t)}{u_t(1 - p_t)}$
Jelinek–Mercer	$P_{JM}(t d) = (1 - \lambda)P_{mle}(t d) + \lambda P(t C)$
Dirichlet smoothing	$P_{Dir}(t d) = \frac{tf_{t,d} + \mu P(t C)}{ d + \mu}$
Naïve Bayes	$c_{\text{map}} = \arg \max_c P(c) \prod_t P(t c)^{tf_{t,d}}; \hat{P}(t c) = \frac{T_{ct} + 1}{\sum_{t'} T_{ct'} + V }$
PageRank	$PR(p) = \frac{1 - d}{N} + d \sum_{q \rightarrow p} \frac{PR(q)}{L(q)}$
Edit distance	$D[i, j] = \min(D[i-1, j] + 1, D[i, j-1] + 1, D[i-1, j-1] + \mathbb{1}[s_1[i] \neq s_2[j]])$
Jaccard coefficient	$J(X, Y) = X \cap Y / X \cup Y $

30 Comprehensive Review Problems

Review Problem 8.1 – Boolean Merge Order (Week 1)

Document frequencies: $df(\text{apple}) = 800$, $df(\text{banana}) = 50$, $df(\text{cherry}) = 200$. For **apple AND banana AND cherry**, give the optimal merge order and an upper bound on the total number of comparisons.

Solution. Order by increasing df : **banana** (50), **cherry** (200), **apple** (800). Merge **banana** and **cherry** first (cost $\leq 50 + 200 = 250$, result ≤ 50 items), then merge that result with **apple** (cost $\leq 50 + 800 = 850$). Total ≤ 1100 comparisons – far less than merging **apple** and **cherry** first ($\leq 800 + 200 = 1000$ for that single step alone).

Review Problem 8.2 – Edit Distance (Week 2)

Compute the edit distance between **pearl** and **earl**.

Solution. **earl** is obtained from **pearl** by deleting the leading **p**. Edit distance = 1.

Review Problem 8.3 – tf-idf and Cosine Similarity (Week 3)

$N = 4$. $df(x) = 4$ (so $idf_x = \log_2(4/4) = 0$), $df(y) = 2$ ($idf_y = \log_2(4/2) = 1$), $df(z) = 1$ ($idf_z = \log_2(4/1) = 2$). Document E has $tf(x) = 6, tf(y) = 3, tf(z) = 2$. Query = “ $y\ z$ ”, giving $Q = (0, 1, 2)$ in idf -weighted form. Compute $\cos(Q, E)$.

Solution. $E = (6 \times 0, 3 \times 1, 2 \times 2) = (0, 3, 4)$. $\|E\| = \sqrt{9 + 16} = 5$. $\|Q\| = \sqrt{1 + 4} = \sqrt{5} \approx 2.236$. $Q \cdot E = 1(3) + 2(4) = 11$.

$$\cos(Q, E) = \frac{11}{5 \times 2.236} = \frac{11}{11.18} \approx 0.984$$

Review Problem 8.4 – BM25 (Week 3)

$N = 5$, $df(t) = 1$ so $idf(t) = \ln \frac{5-1+0.5}{1+0.5} = \ln 3 \approx 1.0986$. $\text{avgdl} = 20$, $k_1 = 1.2$, $b = 0.75$. Document D has $|D| = 20$ (exactly average) and $f(t, D) = 2$. Compute $\text{score}(D, \{t\})$.

Solution. Since $|D| = \text{avgdl}$, the length factor is exactly $1 - b + b(1) = 1$, so the denominator is simply $f(t, D) + k_1 = 2 + 1.2 = 3.2$.

$$\text{score} = 1.0986 \times \frac{2(2.2)}{3.2} = 1.0986 \times \frac{4.4}{3.2} = 1.0986 \times 1.375 \approx 1.511$$

Review Problem 8.5 – Precision, Recall, F_1 (Week 4)

15 documents are relevant to a query. A system retrieves 10 documents, of which 7 are relevant. Compute P , R , F_1 .

Solution. $TP = 7$, $FP = 3$, $FN = 8$. $P = 7/10 = 0.7$, $R = 7/15 \approx 0.467$.

$$F_1 = \frac{2(0.7)(0.467)}{0.7 + 0.467} = \frac{0.653}{1.167} \approx 0.560$$

Review Problem 8.6 – nDCG (Week 4)

Ranked relevance $rel = [1, 2, 0, 3]$. Compute DCG_4 , $IDCG_4$, $nDCG_4$.

Solution. $DCG_4 = 1 + \frac{2}{\log_2 2} + \frac{0}{\log_2 3} + \frac{3}{\log_2 4} = 1 + 2 + 0 + 1.5 = 4.5$.

Ideal order $[3, 2, 1, 0]$: $IDCG_4 = 3 + \frac{2}{1} + \frac{1}{1.585} + \frac{0}{2} = 3 + 2 + 0.631 + 0 = 5.631$.

$$nDCG_4 = \frac{4.5}{5.631} \approx 0.799$$

Review Problem 8.7 – Rocchio (Week 4)

$\vec{q}_0 = (1, 1)$, one relevant document $\vec{d} = (2, 0)$, one non-relevant document $\vec{d}' = (0, 2)$. Using $\alpha = 1$, $\beta = 0.75$, $\gamma = 0.25$, compute $\vec{q}_m$.

Solution.

$$\vec{q}_m = (1, 1) + 0.75(2, 0) - 0.25(0, 2) = (1, 1) + (1.5, 0) - (0, 0.5) = (2.5, 0.5)$$

Review Problem 8.8 – Naïve Bayes (Week 6)

Training: **Fruit** docs – “sweet red apple”, “sweet orange citrus”; **Vegetable** docs – “green leafy spinach”, “green bitter kale”. Test: “sweet green apple”. $|V| = 10$. Classify using add-one smoothing.

Solution. $N = 4$, $N_{\text{Fruit}} = N_{\text{Veg}} = 2 \Rightarrow \hat{P}(\text{Fruit}) = \hat{P}(\text{Veg}) = 0.5$. Fruit counts (total 6): sweet= 2, apple= 1, green= 0. Vegetable counts (total 6): green= 2, sweet= 0, apple= 0.

$$\hat{P}(\text{sweet} \mid F) = \frac{3}{16}, \quad \hat{P}(\text{green} \mid F) = \frac{1}{16}, \quad \hat{P}(\text{apple} \mid F) = \frac{2}{16}$$

$$\hat{P}(\text{sweet} \mid V) = \frac{1}{16}, \quad \hat{P}(\text{green} \mid V) = \frac{3}{16}, \quad \hat{P}(\text{apple} \mid V) = \frac{1}{16}$$

$$\text{Score}(F) = 0.5 \times \frac{3}{16} \times \frac{1}{16} \times \frac{2}{16} \approx 7.32 \times 10^{-4}, \quad \text{Score}(V) = 0.5 \times \frac{1}{16} \times \frac{3}{16} \times \frac{1}{16} \approx 3.66 \times 10^{-4}$$

$\text{Score}(F)$ is exactly twice $\text{Score}(V)$ (both share the factor $\frac{3}{16} \times \frac{1}{16}$; the remaining factors are $\frac{2}{16}$ vs. $\frac{1}{16}$) $\Rightarrow$ classify as **Fruit**.

Review Problem 8.9 – PageRank on a Symmetric Cycle (Week 7)

Three pages $P \rightarrow Q \rightarrow R \rightarrow P$ (a pure cycle, each page has exactly one outlink and one inlink). Using $d = 0.85$, starting from $PR_0 = \frac{1}{3}$ for all three pages, compute PR_1 for each page. What do you notice, and why does it happen?

Solution. Teleport term = $\frac{0.15}{3} = 0.05$. Every page has exactly one inlink from a page with rank $\frac{1}{3}$ and out-degree 1:

$$PR_1(P) = 0.05 + 0.85 \times \frac{1}{3} = 0.05 + 0.2833 = 0.3333 = \frac{1}{3}$$

By the same computation, $PR_1(Q) = PR_1(R) = \frac{1}{3}$ too – **the distribution does not change**. This is because a symmetric cycle is already at its stationary distribution: every page has exactly one inlink carrying exactly the same rank as every other page, so the uniform distribution is a fixed point of Key Formula 7.1.

Appendix

Further Reading

Primary text (all weeks): Manning, C. D., Raghavan, P., & Schütze, H. (2009). *An Introduction to Information Retrieval*. Cambridge University Press. Freely available at <http://www.informationretrieval.org/>.

- **Week 1:** Waller, W. G., & Kraft, D. H. (1979). A mathematical model of a weighted Boolean retrieval system. *Information Processing & Management*, 15(5), 235–245.
- **Week 3:** Lee, D. L., Chuang, H., & Seamons, K. (1997). Document ranking and the vector-space model. *IEEE Software*, 14(2), 67–75. Ramos, J. (2003). Using tf-idf to determine word relevance in document queries. *Proceedings of the First Instructional Conference on Machine Learning*, 242(1), 29–48.
- **Week 4:** Bama, S. S., Ahmed, M. I., & Saravanan, A. (2015). A survey on performance evaluation measures for information retrieval systems. *International Research Journal of Engineering and Technology*, 2(2), 1015–1020. Lavrenko, V., & Croft, W. B. (2017). Relevance-based language models. *ACM SIGIR Forum*, 51(2), 260–267.
- **Week 7:** Bharat, K., & Broder, A. (1998). A technique for measuring the relative size and overlap of public web search engines. *Computer Networks and ISDN Systems*, 30(1–7), 379–388. (The source of the index-overlap technique in Worked Example 7.1.)

Notation Glossary

Symbol	Meaning
N	Number of documents in the collection
V	Vocabulary (set of distinct terms); $ V $ its size
df_t	Document frequency of term t (# documents containing t)
$tf_{t,d}$	Term frequency of t in document d
idf_t	Inverse document frequency of t
$\vec{d}, \vec{q}$	tf-idf vector representations of a document / query
R	(Week 4/5) total number of documents relevant to a query
$P(c)$	Prior probability of class c
$P(t c)$	Probability of term t given class/model c
$PR(p)$	PageRank score of page p
$L(q)$	Out-degree (number of outgoing links) of page q
d	(Week 7 only) PageRank damping factor